Hive Assignment - 1

Adarsh Hegde

1NT18IS011

ISE –A

- Create database called customer
- Create table bank(tid, bname, cname, amount)
- Insert values using insert statement
- Import the values using csv
- Query operation, group by, aggregate, order by and having clause ,views
- Create one more table
- Experiment with join operation ,Alter operations

Create table bank(tid, bname, cname, amount).

Insert values using insert statement.

```
                          hdoop@ubuntu: ~/apache-hive-3.1.2-bin
File  Edit  View  Search  Terminal  Help
Time taken: 0.256 seconds
hive> insert into bank values(501,"Axis","Adarsh",30000.00),(502,"SBI","Narayan",50000.00);
Query ID = hdoop_20210702024255_dce512eb-168e-4602-85be-a00178ef73e8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0008, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0008/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 02:43:11,847 Stage-1 map = 0%,  reduce = 0%
2021-07-02 02:43:31,643 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 20.51 sec
2021-07-02 02:43:41,195 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 25.51 sec
MapReduce Total cumulative CPU time: 25 seconds 510 msec
Ended Job = job_1625216497416_0008
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/bank/.hive-staging_hive_2021-07-02_02-42-55_164_241469877975988
3053-1/-ext-10000
Loading data to table customer.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 25.51 sec   HDFS Read: 19755 HDFS Write: 391 SUCCESS
Total MapReduce CPU Time Spent: 25 seconds 510 msec
OK
Time taken: 48.89 seconds
hive> insert into bank values(503,"Canara","Venkat",40000.00),(504,"KDCC","Vijay",25000.00),(505,"YesBank","Naveen",45000.00);
Query ID = hdoop_20210702024621_aac46b1e-2fc6-4e1c-9bea-60d8648547dc
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
```

```
                          hdoop@ubuntu: ~/apache-hive-3.1.2-bin
File  Edit  View  Search  Terminal  Help
Query ID = hdoop_20210702024621_aac46b1e-2fc6-4e1c-9bea-60d8648547dc
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0009, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0009/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0009
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 02:46:30,791 Stage-1 map = 0%,  reduce = 0%
2021-07-02 02:46:37,034 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.5 sec
2021-07-02 02:46:41,127 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.39 sec
MapReduce Total cumulative CPU time: 4 seconds 390 msec
Ended Job = job_1625216497416_0009
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/bank/.hive-staging_hive_2021-07-02_02-46-21_695_643472262717298
29-1/-ext-10000
Loading data to table customer.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.39 sec   HDFS Read: 19932 HDFS Write: 467 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 390 msec
OK
Time taken: 21.709 seconds
hive> select * from bank;
OK
501     Axis    Adarsh  30000.0
502     SBI     Narayan 50000.0
503     Canara  Venkat  40000.0
504     KDCC    Vijay   25000.0
505     YesBank Naveen  45000.0
Time taken: 0.184 seconds, Fetched: 5 row(s)
hive> 
```

Import the values using csv

```
                                    hdoop@ubuntu: ~/apache-hive-3.1.2-bin
File  Edit  View  Search  Terminal  Help
        at org.apache.hadoop.hive.ql.Driver.compileAndRespond(Driver.java:1773)
        at org.apache.hadoop.hive.ql.Driver.compileAndRespond(Driver.java:1768)
        at org.apache.hadoop.hive.ql.reexec.ReExecDriver.compileAndRespond(ReExecDriver.java:126)
        at org.apache.hadoop.hive.ql.reexec.ReExecDriver.run(ReExecDriver.java:214)
        at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:239)
        at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:188)
        at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:402)
        at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:821)
        at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:759)
        at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:683)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
        at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:498)
        at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
        at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
FAILED: ParseException line 1:7 mismatched input 'data' expecting FROM near 'import' in import statement
hive> load data loacl inpath '/home/hdoop/customer.csv' into table bank;
FAILED: ParseException line 1:10 extraneous input 'loacl' expecting INPATH near '<EOF>'
hive> load data local inpath '/home/hdoop/customer.csv' into table bank;
Loading data to table customer.bank
OK
Time taken: 0.624 seconds
hive> select * from bank;
OK
501     Axis    Adarsh  30000.0
502     SBI     Narayan 50000.0
503     Canara  Venkat  40000.0
504     KDCC    Vijay   25000.0
505     YesBank Naveen  45000.0
506     BOB     Chaitra 35000.0
507     Urban   Nitya   15000.0
508     PNB     Gagan   60000.0
Time taken: 0.1 seconds, Fetched: 8 row(s)
hive> select cname,amount from bank where amount>=30000.00 and group by cname,amount;
NoViableAltException(144@[])
        at org.apache.hadoop.hive.ql.parse.HiveParser_IdentifiersParser.precedenceSimilarExpression(HiveParser_IdentifiersParser.java:9027)
        at org.apache.hadoop.hive.ql.parse.HiveParser_IdentifiersParser.precedenceEqualExpression(HiveParser_IdentifiersParser.java:10254)
```

Query using group by clause

```
                                    hdoop@ubuntu: ~/apache-hive-3.1.2-bin
File  Edit  View  Search  Terminal  Help
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
        at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:498)
        at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
        at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
FAILED: ParseException line 1:57 cannot recognize input near 'group' 'by' 'cname' in expression specification
hive> select cname,amount from bank where amount>=30000.00 group by cname,amount;
Query ID = hdoop_20210702030214_f28b19a6-975f-4f68-9443-da6bcaaa54f7
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0010, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0010/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0010
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:02:23,837 Stage-1 map = 0%,  reduce = 0%
2021-07-02 03:02:39,776 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 15.99 sec
2021-07-02 03:02:47,187 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 17.86 sec
MapReduce Total cumulative CPU time: 17 seconds 860 msec
Ended Job = job_1625216497416_0010
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 17.86 sec   HDFS Read: 13145 HDFS Write: 250 SUCCESS
Total MapReduce CPU Time Spent: 17 seconds 860 msec
OK
Adarsh  30000.0
Chaitra 35000.0
Gagan   60000.0
Narayan 50000.0
Naveen  45000.0
Venkat  40000.0
Time taken: 34.357 seconds, Fetched: 6 row(s)
hive>
```

Query using order by clause

```
File  Edit  View  Search  Terminal  Help
505      Naveen   45000.0
503      Venkat   40000.0
506      Chaitra  35000.0
501      Adarsh   30000.0
504      Vijay    25000.0
507      Nitya    15000.0
Time taken: 19.21 seconds, Fetched: 8 row(s)
hive> select tid,cname,amount from bank where amount>25000 order by amount DESC;
Query ID = hdoop_20210702031350_9a1731e7-f490-4314-8e0a-a86facda75eb
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0013, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0013/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0013
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:13:58,209 Stage-1 map = 0%,   reduce = 0%
2021-07-02 03:14:03,379 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 2.48 sec
2021-07-02 03:14:07,463 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.27 sec
MapReduce Total cumulative CPU time: 4 seconds 270 msec
Ended Job = job_1625216497416_0013
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.27 sec   HDFS Read: 12587 HDFS Write: 274 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 270 msec
OK
508      Gagan    60000.0
502      Narayan  50000.0
505      Naveen   45000.0
503      Venkat   40000.0
506      Chaitra  35000.0
501      Adarsh   30000.0
Time taken: 18.766 seconds, Fetched: 6 row(s)
hive>
```

Query using having clause

```
File  Edit  View  Search  Terminal  Help
506      Chaitra 35000.0
501      Adarsh  30000.0
Time taken: 18.766 seconds, Fetched: 6 row(s)
hive> select tid,cname,amount from bank group by tid having amount>40000;
FAILED: SemanticException [Error 10025]: Line 1:54 Expression not in GROUP BY key '40000'
hive> select tid,cname,sum(amount) from bank group by tid having sum(amount)>40000;
FAILED: SemanticException [Error 10025]: Line 1:11 Expression not in GROUP BY key 'cname'
hive> select tid,cname,sum(amount) from bank group by tid,cname having sum(amount)>40000;
Query ID = hdoop_20210702033442_c459c41d-5e85-4014-8250-a0963096ecc1
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0014, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0014/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:35:17,556 Stage-1 map = 0%,   reduce = 0%
2021-07-02 03:35:43,198 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 10.09 sec
2021-07-02 03:35:57,917 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 13.42 sec
MapReduce Total cumulative CPU time: 13 seconds 420 msec
Ended Job = job_1625216497416_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 13.42 sec   HDFS Read: 14758 HDFS Write: 180 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 420 msec
OK
502      Narayan 50000.0
505      Naveen  45000.0
508      Gagan   60000.0
Time taken: 77.601 seconds, Fetched: 3 row(s)
hive>
```

Query on aggregate functions

```
                              hdoop@ubuntu: ~/apache-hive-3.1.2-bin                      ⊝ ⊡ ⊗
 File  Edit  View  Search  Terminal  Help
Ended Job = job_1625216497416_0010
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 17.86 sec   HDFS Read: 13145 HDFS Write: 250 SUCCESS
Total MapReduce CPU Time Spent: 17 seconds 860 msec
OK
Adarsh  30000.0
Chaitra 35000.0
Gagan   60000.0
Narayan 50000.0
Naveen  45000.0
Venkat  40000.0
Time taken: 34.357 seconds, Fetched: 6 row(s)
hive> select max(amount) as MaxAmt,min(amount)as MinAmt,avg(amount) as AvgAmt from bank;
Query ID = hdoop_20210702030758_1479b2c7-bf26-4c81-96e7-a8d279a60d65
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0011, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0011/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0011
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:08:07,767 Stage-1 map = 0%,  reduce = 0%
2021-07-02 03:08:12,191 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.6 sec
2021-07-02 03:08:19,512 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 4.2 sec
MapReduce Total cumulative CPU time: 4 seconds 200 msec
Ended Job = job_1625216497416_0011
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 4.2 sec   HDFS Read: 17227 HDFS Write: 123 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 200 msec
OK
60000.0 15000.0 37500.0
Time taken: 22.763 seconds, Fetched: 1 row(s)
hive> ▯
```

Construction of view

```
                              hdoop@ubuntu: ~/apache-hive-3.1.2-bin                      ⊝ ⊡ ⊗
 File  Edit  View  Search  Terminal  Help
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0014, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0014/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:35:17,556 Stage-1 map = 0%,  reduce = 0%
2021-07-02 03:35:43,198 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 10.09 sec
2021-07-02 03:35:57,917 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 13.42 sec
MapReduce Total cumulative CPU time: 13 seconds 420 msec
Ended Job = job_1625216497416_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 13.42 sec   HDFS Read: 14758 HDFS Write: 180 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 420 msec
OK
502     Narayan 50000.0
505     Naveen  45000.0
508     Gagan   60000.0
Time taken: 77.601 seconds, Fetched: 3 row(s)
hive> create view Cust_Transac as select tid,cname from bank;
OK
Time taken: 0.466 seconds
hive> select * from Cust_Transac;
OK
501     Adarsh
502     Narayan
503     Venkat
504     Vijay
505     Naveen
506     Chaitra
507     Nitya
508     Gagan
Time taken: 0.139 seconds, Fetched: 8 row(s)
hive> show databases;
```

Creating one more table called customer(cust_id int,tid int,loc string)

And importing values from csv file


```
Time taken: 0.162 seconds, Fetched: 3 row(s)
hive> use customer;
OK
Time taken: 0.074 seconds
hive> show tables;
OK
bank
cust_transac
Time taken: 0.089 seconds, Fetched: 2 row(s)
hive> select * from bank;
OK
501     Axis    Adarsh  30000.0
502     SBI     Narayan 50000.0
503     Canara  Venkat  40000.0
504     KDCC    Vijay   25000.0
505     YesBank Naveen  45000.0
506     BOB     Chaitra 35000.0
507     Urban   Nitya   15000.0
508     PNB     Gagan   60000.0
Time taken: 0.147 seconds, Fetched: 8 row(s)
hive> create table customer(cust_id int,tid int,loc string)row format delimited fields terminated by ",";
OK
Time taken: 0.152 seconds
hive> load data local inpath'/home/hdoop/customer2.csv' into table customer;
Loading data to table customer.customer
OK
Time taken: 0.143 seconds
hive> select * from customer;
OK
101     501     Bangalore
102     503     Mysore
103     502     Hubli
104     505     Dharwad
Time taken: 0.302 seconds, Fetched: 4 row(s)
hive>
```

Experimenting on join operation


```
OK
Time taken: 0.143 seconds
hive> select * from customer;
OK
101     501     Bangalore
102     503     Mysore
103     502     Hubli
104     505     Dharwad
Time taken: 0.302 seconds, Fetched: 4 row(s)
hive> select cname,loc from bank b,customer c where b,tid=c.tid;
FAILED: ParseException line 1:47 missing EOF at ',' near 'b'
hive> select cname,loc from bank b,customer c where b.tid=c.tid;
Query ID = hdoop_20210702035350_e0b925e3-7db1-45fa-9947-8fc04c0cb532
Total jobs = 1
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625216497416_0015, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0015/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0015
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-02 03:54:21,805 Stage-3 map = 0%,  reduce = 0%
2021-07-02 03:54:40,467 Stage-3 map = 100%,  reduce = 0%, Cumulative CPU 7.57 sec
MapReduce Total cumulative CPU time: 7 seconds 570 msec
Ended Job = job_1625216497416_0015
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 7.57 sec   HDFS Read: 9501 HDFS Write: 195 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 570 msec
OK
Adarsh  Bangalore
Narayan Hubli
Venkat  Mysore
Naveen  Dharwad
Time taken: 51.363 seconds, Fetched: 4 row(s)
hive>
```

Experimenting on alter operation

```
File  Edit  View  Search  Terminal  Help
Venkat   Mysore
Naveen   Dharwad
Time taken: 51.363 seconds, Fetched: 4 row(s)
hive> alter table customer rename to cust;
OK
Time taken: 0.237 seconds
hive> show tables;
OK
bank
cust
cust_transac
Time taken: 0.02 seconds, Fetched: 3 row(s)
hive> alter table customer add columns(phno int);
FAILED: SemanticException [Error 10001]: Table not found customer.customer
hive> alter table cust add columns(phno int);
OK
Time taken: 0.371 seconds
hive> desc cust;
OK
cust_id                 int
tid                     int
loc                     string
phno                    int
Time taken: 0.032 seconds, Fetched: 4 row(s)
hive> insert into cust values(207,103,"Sirsi",9998877665);
Query ID = hdoop_20210702040140_48dc6854-6c6f-4a39-9523-f1a5fee85be8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
```

```
File  Edit  View  Search  Terminal  Help
hive> insert into cust values(207,103,"Sirsi",9998877665);
Query ID = hdoop_20210702040140_48dc6854-6c6f-4a39-9523-f1a5fee85be8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0016, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0016/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0016
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 04:01:52,191 Stage-1 map = 0%,  reduce = 0%
2021-07-02 04:02:20,244 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 11.34 sec
2021-07-02 04:02:33,604 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 13.95 sec
MapReduce Total cumulative CPU time: 13 seconds 950 msec
Ended Job = job_1625216497416_0016
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/cust/.hive-staging_hive_2021-07-02_04-01-40_904_200087144685990
7234-1/-ext-10000
Loading data to table customer.cust
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 13.95 sec   HDFS Read: 17826 HDFS Write: 344 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 950 msec
OK
Time taken: 55.071 seconds
hive> select * from cust;
OK
207     103     Sirsi   1408943073
101     501     Bangalore       NULL
102     503     Mysore  NULL
```

```
                              hdoop@ubuntu: ~/apache-hive-3.1.2-bin

File  Edit  View  Search  Terminal  Help
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0016, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0016/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job  -kill job_1625216497416_0016
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 04:01:52,191 Stage-1 map = 0%,  reduce = 0%
2021-07-02 04:02:20,244 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 11.34 sec
2021-07-02 04:02:33,604 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 13.95 sec
MapReduce Total cumulative CPU time: 13 seconds 950 msec
Ended Job = job_1625216497416_0016
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/cust/.hive-staging_hive_2021-07-02_04-01-40_904_200087144685990
7234-1/-ext-10000
Loading data to table customer.cust
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 13.95 sec   HDFS Read: 17826 HDFS Write: 344 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 950 msec
OK
Time taken: 55.071 seconds
hive> select * from cust;
OK
207     103     Sirsi   1408943073
101     501     Bangalore       NULL
102     503     Mysore  NULL
103     502     Hubli   NULL
104     505     Dharwad NULL
Time taken: 0.175 seconds, Fetched: 5 row(s)
hive> 
```