

Adarsh Hegde

1NT18IS011

ISE –A

Hive Assignment 1

```
hadoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
hive> insert into cust values(207,103,"Sirsi",9998877665);
Query ID = hadoop_20210702040140_48dc6854-6c6f-4a39-9523-f1a5fee85be8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0016, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0016/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0016
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 04:01:52,191 Stage-1 map = 0%, reduce = 0%
2021-07-02 04:02:20,244 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 11.34 sec
2021-07-02 04:02:33,604 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 13.95 sec
MapReduce Total cumulative CPU time: 13 seconds 950 msec
Ended Job = job_1625216497416_0016
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/cust/.hive-staging_hive_2021-07-02_04-01-7234-1/-ext-10000
Loading data to table customer.cust
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 13.95 sec HDFS Read: 17826 HDFS Write: 344 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 950 msec
OK
Time taken: 55.071 seconds
hive> select * from cust;
OK
207      103      Sirsi      1408943073
101      501      Bangalore  NULL
102      503      Mysore    NULL
```

```
hadoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
OK
Time taken: 0.143 seconds
hive> select * from customer;
OK
101      501      Bangalore
102      503      Mysore
103      502      Hubli
104      505      Dharwad
Time taken: 0.302 seconds, Fetched: 4 row(s)
hive> select cname,loc from bank b,customer c where b.tid=c.tid;
FAILED: ParseException line 1:47 missing EOF at ',' near 'b'
hive> select cname,loc from bank b,customer c where b.tid=c.tid;
Query ID = hadoop_20210702035350_e0b925e3-7db1-45fa-9947-8fc04c0cb532
Total jobs = 1
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625216497416_0015, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0015/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0015
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-02 03:54:21,805 Stage-3 map = 0%, reduce = 0%
2021-07-02 03:54:40,467 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 7.57 sec
MapReduce Total cumulative CPU time: 7 seconds 570 msec
Ended Job = job_1625216497416_0015
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 7.57 sec HDFS Read: 9501 HDFS Write: 195 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 570 msec
OK
Adarsh Bangalore
Narayan Hubli
Venkat Mysore
Naveen Dharwad
Time taken: 51.363 seconds, Fetched: 4 row(s)
hive>
```

```
hadoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
Venkat Mysore
Naveen Dharwad
Time taken: 51.363 seconds, Fetched: 4 row(s)
hive> alter table customer rename to cust;
OK
Time taken: 0.237 seconds
hive> show tables;
OK
bank
cust
cust_transac
Time taken: 0.02 seconds, Fetched: 3 row(s)
hive> alter table customer add columns(phno int);
FAILED: SemanticException [Error 10001]: Table not found customer.customer
hive> alter table cust add columns(phno int);
OK
Time taken: 0.371 seconds
hive> desc cust;
OK
cust_id      int
tid          int
loc          string
phno        int
Time taken: 0.032 seconds, Fetched: 4 row(s)
hive> insert into cust values(207,103,"Sirsi",9998877665);
Query ID = hadoop_20210702040140_48dc6854-6c6f-4a39-9523-f1a5fee85be8
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
```

```

hadoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0014, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0014/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:35:17,556 Stage-1 map = 0%, reduce = 0%
2021-07-02 03:35:43,198 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 10.09 sec
2021-07-02 03:35:57,917 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 13.42 sec
MapReduce Total cumulative CPU time: 13 seconds 420 msec
Ended Job = job_1625216497416_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 13.42 sec HDFS Read: 14758 HDFS Write: 180 SUCCESS
Total MapReduce CPU Time Spent: 13 seconds 420 msec
OK
502 Narayan 50000.0
505 Naveen 45000.0
508 Gagan 60000.0
Time taken: 77.601 seconds, Fetched: 3 row(s)
hive> create view Cust_Transac as select tid,cname from bank;
OK
Time taken: 0.466 seconds
hive> select * from Cust_Transac;
OK
501 Adarsh
502 Narayan
503 Venkat
504 Vijay
505 Naveen
506 Chaitra
507 Nitya
508 Gagan
Time taken: 0.139 seconds, Fetched: 8 row(s)
hive> show databases;

```

```

hadoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
Time taken: 0.162 seconds, Fetched: 3 row(s)
hive> use customer;
OK
Time taken: 0.074 seconds
hive> show tables;
OK
bank
cust_transac
Time taken: 0.089 seconds, Fetched: 2 row(s)
hive> select * from bank;
OK
501 Axis Adarsh 30000.0
502 SBI Narayan 50000.0
503 Canara Venkat 40000.0
504 KDCC Vijay 25000.0
505 YesBank Naveen 45000.0
506 BOB Chaitra 35000.0
507 Urban Nitya 15000.0
508 PNB Gagan 60000.0
Time taken: 0.147 seconds, Fetched: 8 row(s)
hive> create table customer(cust_id int,tid int,loc string)row format delimited fields terminated by ",";
OK
Time taken: 0.152 seconds
hive> load data local inpath '/home/hadoop/customer2.csv' into table customer;
Loading data to table customer.customer
OK
Time taken: 0.143 seconds
hive> select * from customer;
OK
101 501 Bangalore
102 503 Mysore
103 502 Hubli
104 505 Dharwad
Time taken: 0.302 seconds, Fetched: 4 row(s)
hive>

```


hdoop@ubuntu: ~/apache-hive-3.1.2-bin

File Edit View Search Terminal Help

```
506      Chaitra 35000.0
501      Adarsh 30000.0
Time taken: 18.766 seconds, Fetched: 6 row(s)
hive> select tid,cname,amount from bank group by tid having amount>40000;
FAILED: SemanticException [Error 10025]: Line 1:54 Expression not in GROUP BY key '40000'
hive> select tid,cname,sum(amount) from bank group by tid having sum(amount)>40000;
FAILED: SemanticException [Error 10025]: Line 1:11 Expression not in GROUP BY key 'cname'
hive> select tid,cname,sum(amount) from bank group by tid,cname having sum(amount)>40000;
Query ID = hdoop_20210702033442_c459c41d-5e85-4014-8250-a0963096ecc1
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0014, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0014
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:35:17,556 Stage-1 map = 0%, reduce = 0%
2021-07-02 03:35:43,198 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 10.09 sec
2021-07-02 03:35:57,917 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 13.42 sec
MapReduce Total cumulative CPU time: 13 seconds 420 msec
Ended Job = job_1625216497416_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 13.42 sec HDFS Read: 14758 HDFS Write: 180
Total MapReduce CPU Time Spent: 13 seconds 420 msec
OK
502      Narayan 50000.0
505      Naveen 45000.0
508      Gagan 60000.0
Time taken: 77.601 seconds, Fetched: 3 row(s)
hive>
```

hdoop@ubuntu: ~/apache-hive-3.1.2-bin

File Edit View Search Terminal Help

```
505      Naveen 45000.0
503      Venkat 40000.0
506      Chaitra 35000.0
501      Adarsh 30000.0
504      Vijay 25000.0
507      Nitya 15000.0
Time taken: 19.21 seconds, Fetched: 8 row(s)
hive> select tid,cname,amount from bank where amount>25000 order by amount DESC;
Query ID = hdoop_20210702031350_9a1731e7-f490-4314-8e0a-a86facda75eb
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0013, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0013
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0013
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 03:13:58,209 Stage-1 map = 0%, reduce = 0%
2021-07-02 03:14:03,379 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.48 sec
2021-07-02 03:14:07,463 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.27 sec
MapReduce Total cumulative CPU time: 4 seconds 270 msec
Ended Job = job_1625216497416_0013
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.27 sec HDFS Read: 12587 HDFS Write: 12587
Total MapReduce CPU Time Spent: 4 seconds 270 msec
OK
508      Gagan 60000.0
502      Narayan 50000.0
505      Naveen 45000.0
503      Venkat 40000.0
506      Chaitra 35000.0
501      Adarsh 30000.0
Time taken: 18.766 seconds, Fetched: 6 row(s)
hive>
```


File Edit View Search Terminal Help

Ended Job = job_1625216497416_0010

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 17.86 sec HDFS Read: 13145 HDFS Write: 1

Total MapReduce CPU Time Spent: 17 seconds 860 msec

OK

Adarsh 30000.0

Chaitra 35000.0

Gagan 60000.0

Narayan 50000.0

Naveen 45000.0

Venkat 40000.0

Time taken: 34.357 seconds, Fetched: 6 row(s)

hive> select max(amount) as MaxAmt,min(amount)as MinAmt,avg(amount) as AvgAmt from bank;

Query ID = hdoop_20210702030758_1479b2c7-bf26-4c81-96e7-a8d279a60d65

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625216497416_0011, Tracking URL = http://ubuntu:8088/proxy/application_1

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0011

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2021-07-02 03:08:07,767 Stage-1 map = 0%, reduce = 0%

2021-07-02 03:08:12,191 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.6 sec

2021-07-02 03:08:19,512 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.2 sec

MapReduce Total cumulative CPU time: 4 seconds 200 msec

Ended Job = job_1625216497416_0011

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.2 sec HDFS Read: 17227 HDFS Write: 12

Total MapReduce CPU Time Spent: 4 seconds 200 msec

OK

60000.0 15000.0 37500.0

Time taken: 22.763 seconds, Fetched: 1 row(s)

hive> □

File Edit View Search Terminal Help

at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)

at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:37)

at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

at java.lang.reflect.Method.invoke(Method.java:498)

at org.apache.hadoop.util.RunJar.run(RunJar.java:323)

at org.apache.hadoop.util.RunJar.main(RunJar.java:236)

FAILED: ParseException line 1:57 cannot recognize input near 'group' 'by' 'cname' in

hive> select cname,amount from bank where amount>=30000.00 group by cname,amount;

Query ID = hdoop_20210702030214_f28b19a6-975f-4f68-9443-da6bcaaa54f7

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625216497416_0010, Tracking URL = http://ubuntu:8088/proxy/appl

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0010

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2021-07-02 03:02:23,837 Stage-1 map = 0%, reduce = 0%

2021-07-02 03:02:39,776 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 15.99 sec

2021-07-02 03:02:47,187 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 17.86 sec

MapReduce Total cumulative CPU time: 17 seconds 860 msec

Ended Job = job_1625216497416_0010

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 17.86 sec HDFS Read: 13145 HDFS

Total MapReduce CPU Time Spent: 17 seconds 860 msec

OK

Adarsh 30000.0

Chaitra 35000.0

Gagan 60000.0

Narayan 50000.0

Naveen 45000.0

Venkat 40000.0

Time taken: 34.357 seconds, Fetched: 6 row(s)

hive> □


```

hdoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
at org.apache.hadoop.hive.ql.Driver.compileAndRespond(Driver.java:1773)
at org.apache.hadoop.hive.ql.Driver.compileAndRespond(Driver.java:1768)
at org.apache.hadoop.hive.ql.reexec.ReExecDriver.compileAndRespond(ReExecDriver.java:12)
at org.apache.hadoop.hive.ql.reexec.ReExecDriver.run(ReExecDriver.java:214)
at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:239)
at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:188)
at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:402)
at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:821)
at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:759)
at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:683)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
FAILED: ParseException line 1:7 mismatched input 'data' expecting FROM near 'import' in import
hive> load data local inpath '/home/hdoop/customer.csv' into table bank;
FAILED: ParseException line 1:10 extraneous input 'load' expecting INPATH near '<EOF>'
hive> load data local inpath '/home/hdoop/customer.csv' into table bank;
Loading data to table customer.bank
OK
Time taken: 0.624 seconds
hive> select * from bank;
OK
501      Axis      Adarsh  30000.0
502      SBI       Narayan 50000.0
503      Canara    Venkat  40000.0
504      KDCC      Vijay   25000.0
505      YesBank   Naveen 45000.0
506      BOB       Chaitra 35000.0
507      Urban    Nitya   15000.0
508      PNB       Gagan   60000.0
Time taken: 0.1 seconds, Fetched: 8 row(s)
hive> select cname,amount from bank where amount>=30000.00 and group by cname,amount;
NoViableAltException(144@[])
at org.apache.hadoop.hive.ql.parse.HiveParser_IdentifiersParser.precedenceSimilarExpress
at org.apache.hadoop.hive.ql.parse.HiveParser_IdentifiersParser.precedenceEqualExpressi

```

```

hdoop@ubuntu: ~/apache-hive-3.1.2-bin
File Edit View Search Terminal Help
Query ID = hdoop_20210702024621_aac46b1e-2fc6-4e1c-9bea-60d8648547dc
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625216497416_0009, Tracking URL = http://ubuntu:8088/proxy/application
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0009
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-02 02:46:30,791 Stage-1 map = 0%, reduce = 0%
2021-07-02 02:46:37,034 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.5 sec
2021-07-02 02:46:41,127 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.39 sec
MapReduce Total cumulative CPU time: 4 seconds 390 msec
Ended Job = job_1625216497416_0009
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/bank/.hive
29-1/-ext-10000
Loading data to table customer.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.39 sec HDFS Read: 19932 HDFS Write:
Total MapReduce CPU Time Spent: 4 seconds 390 msec
OK
Time taken: 21.709 seconds
hive> select * from bank;
OK
501      Axis      Adarsh  30000.0
502      SBI       Narayan 50000.0
503      Canara    Venkat  40000.0
504      KDCC      Vijay   25000.0
505      YesBank   Naveen 45000.0
Time taken: 0.184 seconds, Fetched: 5 row(s)
hive> 

```

File Edit View Search Terminal Help

Time taken: 0.256 seconds

hive> insert into bank values(501,"Axis","Adarsh",30000.00),(502,"SBI","Narayan",50000.00);

Query ID = hdoop_20210702024255_dce512eb-168e-4602-85be-a00178ef73e8

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1625216497416_0008, Tracking URL = http://ubuntu:8088/proxy/application_1625216497416_0008

Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625216497416_0008

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2021-07-02 02:43:11,847 Stage-1 map = 0%, reduce = 0%

2021-07-02 02:43:31,643 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 20.51 sec

2021-07-02 02:43:41,195 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 25.51 sec

MapReduce Total cumulative CPU time: 25 seconds 510 msec

Ended Job = job_1625216497416_0008

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/bank/.hive-staging_h

3053-1/-ext-10000

Loading data to table customer.bank

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 25.51 sec HDFS Read: 19755 HDFS Write: 391 SUC

Total MapReduce CPU Time Spent: 25 seconds 510 msec

OK

Time taken: 48.89 seconds

hive> insert into bank values(503,"Canara","Venkat",40000.00),(504,"KDCC","Vijay",25000.00),(505,"Y

Query ID = hdoop_20210702024621_aac46b1e-2fc6-4e1c-9bea-60d8648547dc

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>